

k-NN and range search with kd-trees

Dominika Kubániová

November 11, 2019

kd-tree

- ▶ Binary tree
- ▶ Data points $\in \mathbb{R}^k$ (my case k is max 4)
- ▶ Root wraps data into hyperrectangle (cell)
- ▶ Hyperrectangle iteratively splitted by orthogonal planes

Non-optimized structures

```
1 struct kd_node {  
2     kd_node * left, * right;  
3     byte split_dim;  
4     float split_value;  
5     bool is_leaf;  
6     float * p;  
7 }
```

Note: Some implementations store points in nodes and leaves, some only in leaves

Search problems

- ▶ k-NN search
- ▶ Range search (rectangular and spherical queries)

Construction of a kd-tree

```
1 fun construct_kdtree(dataset of size N in dimension k):  
2   if dataset size < M:  
3     return nullptr  
4   kd_node * node  
5   [split_dim, split_value] = choose_split()  
6   // ...  
7   data_left = {p from dataset | p[dim] < value}  
8   data_right = {p from dataset | p[dim] > value}  
9   node->left = construct_kdtree(data_left)  
10  node->right = construct_kdtree(data_right)  
11  return node  
12
```

Data should be sorted in every dimension

How to choose splitting dimension?

- ▶ Does not affect correctness of the algorithm
- ▶ Can increase complexity of the search
- ▶ Balanced tree
- ▶ We want ratio of the longest to shortest side to be bounded
- ▶ The number of visited leaves in search should be small
- ▶ Bad dataset distribution can ruin everything



Dimension with maximal variance

- ▶ Splitting value = median
- ▶ Pros:
 - ▶ $O(N)$ leaves and $O(\log N)$ depth
 - ▶ Probably every next chosen dimension will be different (not true when points have bad distribution)
- ▶ Cons:
 - ▶ Unbounded side length ratio (median does not need to be in the middle of the hyperrectangle)

Midpoint split

- ▶ Splitting value = half of the longest side of the hyperrectangle
- ▶ Balanced
- ▶ Pros:
 - ▶ Side length ratio bounded
- ▶ Cons:
 - ▶ Possibly a lot of empty nodes

Max variance and Midpoint




Source [2]

Large number of leaves is visited in k-NN search!


Sliding midpoint

- ▶ First step as midpoint
- ▶ If points lie on both sides of this split \rightarrow continue
- ▶ Else \rightarrow slide towards the closest point creating one NON-EMPTY leaf from it
- ▶ Pros:
 - ▶ $O(N)$ leaves
- ▶ Cons:
 - ▶ What about its side length ratio? Is it balanced?



Source: [2]

It is observed that the depth of the tree is not that dominant factor than the number of visited leaves. Example:



Source: [2]

k-NN search (with PQ)

```
1 Point query
2 float knn_dist = INF
3 PQ kNNq
4 PQ q
5
6 fun process_leaf(n):
7     for p in n->points:
8         if (kNNq.size) < k:
9             kNNq.push(p)
10        else if dist(p, query) < knn_dist:
11            kNNq.push(p)
12            if kNNq.size() > k:
13                kNNq.pop_last()
14            knn_dist = dist(kNNq.last, query)
15
```

k-NN search (with PQ) cont.

```
1 q.push([root, 0.0])
2 while !q.empty():
3     [n, dist] = q.pop()
4     if dis < knn_dist:
5         while (!n->is_leaf):
6             off_split_plane = query[n->split_dim] - n->
split_value
7             if off_split_plane < 0:
8                 off_box_boundary = query[n->split_dim] - n->
bound
9                 if off_box_boundary > 0:
10                     off_box_boundary = 0
11                     dist_right = dist - off_box_boundary^2 +
off_split_plane^2
12                     q.push([n->right, dist_right])
13                     n = n->left
14             else:
15                 // analogously
16                 process_leaf(n)
17
```

Range search (spherical query)

```
1 Point query
2 float r_sq
3 float offs[k] // initialized 0
4 vector<Point> in_range
5 stack s
6
7 fun process_leaf(n):
8     for p in n->points:
9         if dist(p, query) < r_sq:
10             in_range.push(p)
11
```

Range search (spherical query) cont.

```
1 s.push([root, dist, offs])
2 while !s.empty():
3     [n, dist, off] = s.pop()
4     if n->is_leaf:
5         process_leaf(n)
6     else:
7         old_off_split_plane = off[n->dim]
8         new_off_split_plane = query[n->dim] - n->
split_value
9         if new_off_split_plane < 0:
10             new_dist = dist - old_off_split_plane^2 +
new_off_split_plane^2
11             if new_dist < r_sq:
12                 new_off = off
13                 new_off[n->dim] = new_off_split_plane
14                 s.push([n->right, dist, new_off])
15                 s.push([n->left, dist, off])
16             else:
17                 // analogously
18
```

Range search (rectangular query)

```
1 Point query
2 float size[k] // only halves stored
3 float offs[k] // initialized 0
4 vector<Point> in_range
5 stack s
6
7 fun process_leaf(n):
8     for p in n->points:
9         dim = n->dim
10        for i in range(k):
11            if query[dim] - size[dim] < p[dim] < query[dim] +
               size[dim]
12                in_range.push(p)
13            else:
14                break
15
```


Range search (rectangular query) cont.


```
1 s.push(root)
2 while !s.empty():
3     n = s.pop()
4
5     if n->is_leaf:
6         process_leaf(n)
7     else:
8         if n->value < query[n->dim] - size[n->dim]:
9             s.push(n->right)
10
11         else if n->value > query[n->dim] + size[n->dim]:
12             s.push(n->left)
13
14         else:
15             s.push(n->left)
16             s.push(n->right)
17
```

Optimized structures

```
1 struct kd_node {  
2     union {  
3         kd_node * left;  
4         float * p;  
5     }  
6     byte is_leaf_split_dim;  
7     float split_value;  
8 }  
9
```

Other simple heuristics

- Bounds-Overlap-Ball (BOB) test
Effective when number of points stored in a leaf is high



Source: [2]

What next needs to be done

- ▶ Dataset generation (various distributions) in 2D-4D
- ▶ Dataset size: $10^3 - 10^8$
- ▶ Different implementation of priority queue
- ▶ Comparison (mainly with naïve algorithm)

References

- ▶ 1 Moore, Andrew. (2004). An Introductory Tutorial on Kd-Trees.
- ▶ 2 Maneewongvatana, Songrit & Mount, David. (2000). It's Okay to Be Skinny, If Your Friends Are Fat.
- ▶ 3 Sample, Neal & Haines, Matthew & Arnold, Mark & Purcell, Timothy. (2001). Optimizing Search Strategies in k-d Trees.

Thank you for your attention